

# REINFORCEMENT LEARNING (RL) and Q-LEARNING

## 1. Introduction

RL is a method of learning that uses a reward function to guide the agent in the search for the best path to a goal. RL system has the follow components:

- $s$ : **state** describes the observation of the environment.
- $a$ : **action** describes how the agent interacts with the environment.
- $\pi(s)$ : **policy** describes the rule that the agent decides to take an action  $a$  at the state  $s$ .
- $r$ : **reward** describes the reward for the agent after performing an action.
- $\gamma$ : **discount factor** describes how much the agent cares about future rewards.
- $V$ : value function
- $Q$ : (quality) action-value function

## 2. Agent's Learning Objectives:

- Each time step  $t$ , the agent chooses an action based on the current state, he will receive a reward  $r_t$ .
- Along the path, by following the policy  $\pi$ , the total reward  $R$  will be accumulated.

$$R_\pi = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \gamma^3 r_{t+3} + \cdots \gamma^T r_{t+T}$$

where  $\gamma$  is the discount factor, which emphasizes the recent rewards are more important than the future ones.

- The **ultimate goal** is to achieve the maximum accumulated reward  $R$ . However, since the environment is stochastic, the agent will not always get the same reward  $r_t$  for the same action  $a$ . Therefore, the accumulated reward  $R$  is expressed in term of expectation, which is called **Value**.


$$V_\pi(s) = E[r_t + r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \cdots \gamma^T r_{t+T} | \pi(s)]$$

## 3. How to choose the action?

- By the learning objective, we should of course choose the action that maximizes the value function. This is expressed by the Bellman equation:

$$V(s) = \max_a [r(s, a) + \gamma V(s')]$$

In the example below, the agent gets a reward  $r = 1$  when reaching the goal,  $r = -1$  when it gets into fire, and  $r = 0$  otherwise. We use discount factor  $\gamma = 0.9$ .



- We start from near the goal  $s = (3, 3)$ . The value of this state is:  $V(s = (3, 3)) = 1$ , where reward  $r = 1$  and stop (no further state).

- Then, we move 1 step backward to the left  $s = (3, 2)$ . The value of this state is:  $V(s = (3, 2)) = 0 + 0.9 * V(s = (3, 3)) = 0.9$ .
- Then, we continue the steps to all other states.


This procedure is basically the "Dynamic Programing".

- We build the value map starting from the location nearing goal, and extend to all the other location in the map.
- From the Value map above, we can easily set the plan to move the agent.


#### 4. What if the environment is uncertain ?

In the example above, the action is 100% certain, e.g **turn-left** command will definitely move the agent to the left. However, in practice, there are always uncertainties, e.g **turn-left** command will move the agent to the left in 80% cases but 10% cases it can move to the right or forward.

##### Deterministic Search



##### Non-Deterministic Search




To make the problem easier, we requires the process to be an **Markov Process**, that is:

In Markow process, the future state only depends on the current state, not the sequence in the past.

And the Bellman is rewritten in the stochastic form:

$$V(s) = \max_a [r(s, a) + \gamma \sum_{s'} P(s'|s, a) V(s')]$$

#### 5. Q-Learning Intuition



- From the Bellman equation, let we define the quality of an action (left-figure) as:

$$Q(s, a) = r(s, a) + \gamma \sum_{s'} P(s'|s, a) V(s') \Rightarrow V(s) = \max_a [Q(s, a)]$$

The value of the state  $V(s)$  is the maximum quality of an action  $Q(s, a)$ .

- Then substitute  $V(s') = \max_a [Q(s', a)]$  (right-figure), we write the Q-function independent of  $V(s)$ :


$$Q(s, a) = r(s, a) + \gamma \sum_{s'} P(s'|s, a) \max_a [Q(s', a)]$$

- For short notation, we can ignore the "Stochastic Process" and simplify the Q-function as follows.

$$Q(s, a) = r(s, a) + \gamma \max_a [Q(s', a)]$$

## 6. Temporal Difference (TD)

- Theoretically, if  $Q(s, a)$  is known ahead, then every thing is very easy, we just choose the action that maximizes the quality.
- However,  $Q(s, a)$  is just an estimation, which hopefully is getting closer and closer to the true value along the learning path. Therefore, after taking an action, there is always a mismatched from the previous and the re-estimated quality. This amount of mismatch is called **temporal difference (TD)**.




- And we update the Q-function as follows:

$$Q_t(s, a) = Q_{t-1}(s, a) + \alpha TD_t(s, a)$$

Note that, when  $\alpha = 1$ , this reduces to the ideal Bellman equation, i.e. we discard the whole experience learned in the past, and update the Q-function with the latest quality. However, this is not a good idea, because the new estimated value is not necessarily the best value representing the state. For example, it may be a rare event due to the random process.

## 7. Q-Learning Algorithm

**Initialization (First iteration):**

For all couples of states  $s$  and actions  $a$ , the Q-values are initialized to 0.

**Next iterations:**

At each iteration  $t \geq 1$ , we repeat the following steps:

1. We select a random state  $s_t$  from the possible states.
2. From that state, we play a random action  $a_t$ .
3. We reach the next state  $s_{t+1}$  and we get the reward  $R(s_t, a_t)$ .

4. We compute the Temporal Difference  $TD_t(s_t, a_t)$ :

$$TD_t(s_t, a_t) = R(s_t, a_t) + \gamma \max_a (Q(s_{t+1}, a)) - Q(s_t, a_t)$$

5. We update the Q-value by applying the Bellman equation:

$$Q_t(s_t, a_t) = Q_{t-1}(s_t, a_t) + \alpha TD_t(s_t, a_t)$$

- The algorithm is implemented at [q\\_learning.py](#).
- An example of cartpole is shown in [cartpole.py](#).
- An example of mountain car is shown in [mountain\\_car.py](#).