

Twin-Delayed Deep Deterministic Policy Gradient (TD3)

Fujimoto, S., Van Hoof, H., & Meger, D. (2018). Addressing Function Approximation Error in **Actor-Critic** Methods. 35th International Conference on Machine Learning, ICML 2018, 4, 2587–2601.

1. Motivation

- In DDPG, we use a Critic network (q-value) to evaluate and train the Actor network (policy), hence the actor's performance strongly depends on how well Critic network estimates the action-value **q**.
- Due to bootstrap learning mechanism, i.e learning by temporal-difference, the **Critic** is subject to estimation error, namely the overestimation bias and high variance build-up.
- In Double Q-Learning, we know that using an extra Q-network can help reduce this bias. This is the key idea of TD3: We use an extra network Critic, named twin-network, to perform another estimation of **q-value**. By having two Critic networks compete to each other, we can elivate the over-confident estimation error.

2. Twin-Delayed DDPG (TD3)

- Compare two DDPG, the main change is that TD3 uses two Critic networks, namely $Q_1(s, a|\theta_1)$ and $Q_2(s, a|\theta_2)$, which are identical (twin) but with different iniatialized weights θ_1 and θ_2 .
 - Let $Q_{T1}(s', a'|\theta_{T1})$ and $Q_{T2}(s', a'|\theta_{T2})$ be their corresponding target, respectively.
 - Also, let $A(s|\phi)$ be the Actor network, and $A_T(s'|\phi_T)$ be the Actor-target network.
- We update the **q_target** as follows:
 - From the Replayed-Buffer, we sample a mini-batch (s, a, r, s') .
 - We predict the next action for the next state s' using the Actor-target network, disturbed by a small Gaussian noise:

$$a' = A_T(s'|\phi_T) + N(0, \sigma)$$

- We compute two action-values using the twin-Critic Target networks:

$$q_1 = Q_{T1}(s', a'|\theta_{T1}), q_2 = Q_{T2}(s', a'|\theta_{T2})$$

- We compute the **q_target** by

$$q_{target} = r + \gamma \min(q_1, q_2)$$

- Then, we update the Critic Networks by:
 - Evaluate the current action-values using the twin-Critic networks:

$$q_1 = Q_1(s', a'|\theta_1), q_2 = Q_2(s', a'|\theta_2)$$

- Compute the loss to update the Critic-Networks using MSE loss:

$$Loss = \frac{1}{2} [(q_1 - q_{target})^2 + (q_2 - q_{target})^2]$$

- The Actor network $A(s|\phi)$ is trained as done in DDPG, using either one of the Critic, e.g. Q_1 :
 - We predict the action for the current state s using the Actor network:

$$a = A(s|\phi)$$

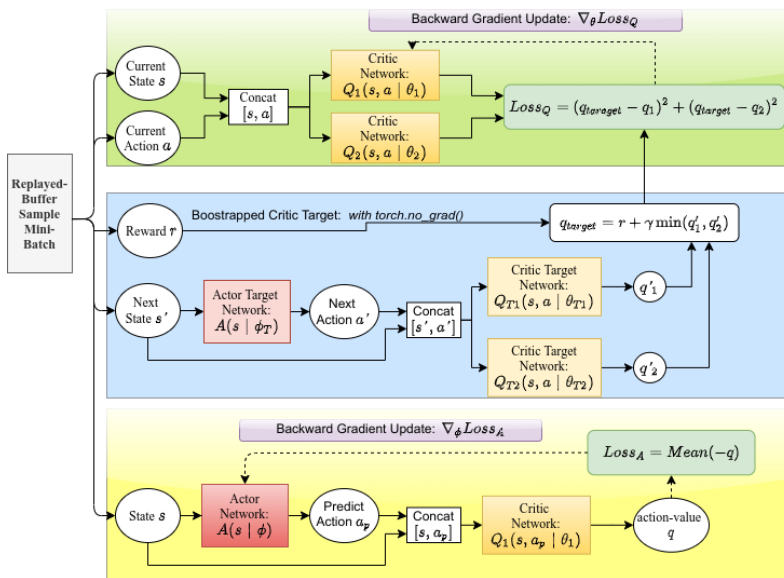
- Loss for Actor network:

$$Loss = Mean(-Q_1(s, a|\theta_1))$$

3. What makes TD3 works?

- As mentioned above, using the Twin-Critic Network helps reducing the over-confident estimation for `q_target` value.
- An important trick is that, we `delay` updating the Actor and the target networks. For every two iterations, we update them **once**.

4. TD3 Pseudo Code



Algorithm 1 TD3

Initialize critic networks $Q_{\theta_1}, Q_{\theta_2}$, and actor network π_ϕ with random parameters θ_1, θ_2, ϕ
 Initialize target networks $\theta'_1 \leftarrow \theta_1, \theta'_2 \leftarrow \theta_2, \phi' \leftarrow \phi$
 Initialize replay buffer \mathcal{B}
for $t = 1$ **to** T **do**
 Select action with exploration noise $a \sim \pi_\phi(s) + \epsilon$, $\epsilon \sim \mathcal{N}(0, \sigma)$ and observe reward r and new state s'
 Store transition tuple (s, a, r, s') in \mathcal{B}

 Sample mini-batch of N transitions (s, a, r, s') from \mathcal{B}
 $\tilde{a} \leftarrow \pi_{\phi'}(s') + \epsilon$, $\epsilon \sim \text{clip}(\mathcal{N}(0, \tilde{\sigma}), -c, c)$
 $y \leftarrow r + \gamma \min_{i=1,2} Q_{\theta'_i}(s', \tilde{a})$
 Update critics $\theta_i \leftarrow \arg\min_{\theta_i} N^{-1} \sum (y - Q_{\theta_i}(s, a))^2$
 if $t \bmod d$ **then**
 Update ϕ by the deterministic policy gradient:
 $\nabla_\phi J(\phi) = N^{-1} \sum \nabla_a Q_{\theta_1}(s, a)|_{a=\pi_\phi(s)} \nabla_\phi \pi_\phi(s)$
 Update target networks:
 $\theta'_i \leftarrow \tau \theta_i + (1 - \tau) \theta'_i$
 $\phi' \leftarrow \tau \phi + (1 - \tau) \phi'$
 end if
end for

To train DDPG agent for Mountain Car Continuous problem, do:

```
python tools/train.py configs/TD3/td3_mountaincar_continuous.py
```

Result after training 260 episodes:

